4A

```cpp
#include <iostream>
#include <cuda_runtime.h>
#include /usr/local/cuda/include/cuda_runtime.h

_global_void addVectors(int* A, int* B, int* C, int n)
{
        int i = blockIdx.x * blockDim.x + threadIdx.x;
        if (i < n)
        {
        C[i] = A[i] + B[i];
        }
}

int main()
{
        int n = 1000000;
        int* A, * B, * C;
        int size = n * sizeof(int);

        // Allocate memory on the host
        cudaMallocHost(&A, size);
        cudaMallocHost(&B, size);
        cudaMallocHost(&C, size);

        // Initialize the vectors
        for (int i = 0; i < n; i++)
        {
        A[i] = i;
        B[i] = i * 2;
        }
        // Allocate memory on the device
        int* dev_A, * dev_B, * dev_C;
        cudaMalloc(&dev_A, size);
        cudaMalloc(&dev_B, size);
        cudaMalloc(&dev_C, size);

        // Copy data from host to device
        cudaMemcpy(dev_A, A, size, cudaMemcpyHostToDevice);
        cudaMemcpy(dev_B, B, size, cudaMemcpyHostToDevice);

        // Launch the kernel
        int blockSize = 256;
        int numBlocks = (n + blockSize - 1) / blockSize;
        // Copy data from device to host
        cudaMemcpy(C, dev_C, size, cudaMemcpyDeviceToHost);

        // Print the results
        for (int i = 0; i < 10; i++)
        {
        cout << C[i] << " ";
```

```
    }
    cout << endl;

    // Free memory
    cudaFree(dev_A);
    cudaFree(dev_B);
    cudaFree(dev_C);
    cudaFreeHost(A);
    cudaFreeHost(B);
    cudaFreeHost(C);

    return 0;
}
```